



João Blasques

AI-Enabled Data Engineer — ~8 years across data & software engineering, 3 of them full-time in data engineering.

📍 Cascais, Portugal ✉️ joaoeduardoblasques@gmail.com 🔗 joaoblasques.com
🐙 github.com/joaoblasques 🔗 linkedin.com/in/joaoblasques

SUMMARY

I build data platforms the way software should be built. Tested transforms. CI/CD. Infrastructure as code. Data-quality gates that fail the build instead of writing a report nobody reads. My software and data background goes back to 2002, at the European Commission's Joint Research Centre, and continued at Universidade Nova de Lisboa (web development, GIS data analysis, BI). I've been doing data engineering full-time since 2023.

The chapter in between was several years repairing wind-turbine blades on ropes, offshore. That was deliberate, hands-on, high-stakes work where a mistake shows up immediately. I brought that same discipline into software: document the trade-offs, check claims against evidence, don't ship what I can't defend.

"AI-enabled" means I run an agent-operated development workflow (Claude Code, CLI-first infrastructure) and I build AI systems too: feature stores, embedding search, LLM-powered pipelines. That work is part of the data platform I build, not a separate thing I do on the side.

EXPERIENCE

Data Engineer — DataStuff.nl (freelance) — Oct 2023 – Present Remote, part-time/freelance. I build governed lakehouses on Databricks and GCP: medallion architecture, dbt transformations and tests, Great Expectations data-quality gates wired into CI, Spark Structured Streaming, Terraform with keyless (WIF) CI. Public showcase of this work: [Vitals](#) (a healthcare lakehouse) and [MBTA](#) (a streaming transit lakehouse).

Blade Repair Technician — various contracts (Muehlhan Wind, GEV Wind Power, Blade defence, Global Wind Service) — Apr 2021 – Dec 2024 Offshore and onshore wind farms in Sweden, Norway, Greece, and Germany, on Vestas and Siemens Gamesa/GE turbines. Leading-edge protection and overlamination repair, uptower upgrades, de-icing and lightning-protection system commissioning, rope-access inspection. The final contract (Global Wind Service, Feb – Dec 2024) overlapped with the DataStuff.nl freelance work above — that overlap was the transition period before data engineering became full-time.

Web Developer / GIS Analyst / Project Manager — Universidade Nova de Lisboa — 2009 – Apr 2013 Lisbon, Portugal. Ran web development projects end to end — project management, analysis, design, implementation, and monitoring. GIS data analysis on spatial datasets, alongside the GeoSpatial Technologies master's at NOVA IMS.

Information Technology Engineer — European Commission, Joint Research Centre — Feb 2002 – Feb 2003 Italy. Web programming, database management, and project implementation.

EDUCATION

- **Master's degree, GeoSpatial Technologies** — NOVA IMS Information Management School (2010 – 2012)
- **Bachelor's degree, Environmental Engineering Technology** — Universidade Nova de Lisboa (1995 – 2001)

CERTIFICATIONS

- **Databricks Fundamentals** — Databricks (issued Jun 2025)
- **Data Engineer in Python** — DataCamp (issued Feb 2025)
- **Associate Data Engineer in SQL** — DataCamp (issued Jan 2025)

SKILLS

Three tiers, each backed by something I've actually shipped and can point to. Full detail with per-skill evidence links: joaoblasques.com/skills. I'm not claiming anything below at a depth the public work doesn't support.

Core — use daily, will defend in depth: Python, SQL · Databricks (Unity Catalog, Delta, Asset Bundles, serverless Jobs, Lakeflow/DLT) · Lakehouse design (medallion, policy-in-gold) · dbt · Data-quality gates (Great Expectations + dbt tests in CI) · CI/CD & IaC (GitHub Actions, Terraform, keyless WIF/OIDC) · Cost discipline (DBU attribution, budget alerts)

Working — shipped with it, still deepening: Spark Structured Streaming (Auto Loader, Kafka) · GCP data stack (BigQuery — external tables, partitioning, clustering, cost tuning · GCS · Pub/Sub · Cloud Run · Cloud Scheduler) · Airflow · Dimensional modelling (star schema) · Feast feature store · pgvector / embeddings · MLflow · Agentic/LLM pipeline patterns · Docker

Familiar — used, not claiming depth: DuckDB & Postgres · Azure/AWS fundamentals · Kubernetes basics · Scala/Java (read comfortably) · Tableau/PowerBI basics

BEYOND WORK

I volunteer time and skills where they're useful — disaster-relief work (2004 Indian Ocean tsunami, Tamil Nadu) and long-term social projects in India and the Philippines.